

Conversational AI Evaluation Report

Executive Summary

This report evaluates a multi-turn travel-planning conversation.

Evaluation Scores

	Metric	Value
0	Goal Completion	True
1	Constraint Consistency	True
2	Adversarial Resistance	True

Failure Types

No failures detected.

Evidence

	Field	Value
0	budget_limit	500
1	final_total_cost	450

Conversation Transcript

	role	content
0	user	Plan me a 3 day trip to Las Vegas from Los Ang...
1	ai-assistant	Okay, let's plan a fantastic 3-day trip to Las...
2	user	Lets's make it a one day trip keeping the same...
3	ai-assistant	Okay, let's plan a fantastic one-day trip with...

Evidently Diagnostics

A detailed data summary report was generated at:

- `outputs/evidently_conversation_report.html`

Conclusion

This evaluation demonstrates:

- Session-level scoring
- Constraint adherence validation
- Multi-turn consistency analysis
- Failure taxonomy generation